

PROFESSIONAL DEVELOPMENT MODULE

AI Tools Landscape for Educators

A curated navigator’s guide to the AI tools relevant to K–12 schools — what each category does, which tools matter, how to evaluate new ones, and the red flags that should stop you in your tracks

Module 1.3 • Track 1: AI Foundations & Literacy • xvalidateai.com

Audience K–12 Teachers	Format Micro-module	Duration 25–30 minutes	Level Intermediate	Delivery Self-paced or facilitated
----------------------------------	-------------------------------	----------------------------------	------------------------------	--

Module Overview

The AI tools available to schools are multiplying rapidly, and the marketing around them ranges from genuinely useful to wildly overstated. Teachers are being asked to evaluate, adopt, or advise on AI tools without a clear framework for doing so. This module gives educators a structured map of the landscape — what the main categories of AI tools are, which specific tools matter most in a school context, how to evaluate any new tool before recommending or adopting it, and what the red flags look like when a tool should not be used with students.

This is not a comprehensive product review. The specific tools landscape changes quickly, and any list will be partially out of date within months. The goal is to give educators a durable framework and evaluation vocabulary that works regardless of which new tools emerge.

IMPORTANT
NOTE

All tools mentioned in this module are included for educational awareness. Inclusion here does not constitute endorsement by xvalidateai.com. Schools must conduct their own vetting using the framework in Section 5 before adopting any AI tool. Tool capabilities, pricing, and data policies change frequently — always verify current terms before use.

Learning Objectives

By the end of this module, participants will be able to:

1	Name and describe the six main categories of AI tools relevant to K–12 education and give at least one example from each
2	Identify the primary use cases, strengths, and risks of the ten most commonly encountered AI tools in schools
3	Apply the SAFE Evaluation Framework to assess any unfamiliar AI tool before using or recommending it

4	Recognise eight specific red-flag patterns that indicate a tool should not be used with students or student data
5	Locate and interpret a tool's privacy policy to determine whether it meets minimum standards for school use
6	Make an evidence-based recommendation to adopt, trial, or decline a specific AI tool for their school context

Section 1: How to Read the AI Tools Landscape

Two principles before looking at any specific tool

The most important conceptual move before evaluating any AI tool is to separate the marketing from the mechanism. The AI tools space is subject to extraordinary levels of hype, and product claims frequently outrun what tools actually deliver. Two principles cut through most of the noise.

Principle 1: Category before product

Understanding what category a tool belongs to tells you more about its fundamental capabilities and risks than any marketing description. A tool that generates text has the same structural strengths and weaknesses as all text-generation tools, regardless of how it is branded or what subject it claims to specialise in.

*Ask: What type of AI is this actually doing?
That answer is more stable and more useful than what the product page says.*

Principle 2: Student data is the highest-stakes decision

Any AI tool that processes, stores, or learns from student data is subject to FERPA, COPPA, and state-level regulations — regardless of whether the vendor acknowledges this. The most common compliance failure in schools is not deliberate wrongdoing: it is a teacher recommending a free AI tool without knowing that it stores student inputs and uses them for training.

Ask: What happens to the data students put into this tool?

The six categories of AI tools in education

Every AI tool a teacher or student is likely to encounter falls into one of these six categories. Understanding the category immediately tells you what the tool can and cannot do, what its failure modes look like, and what the data risks are.

Category	What it does	Core capability	Primary risk
Generative text AI	Generates human-like text in response to prompts	Writing assistance, explanation, feedback drafts, lesson planning, Q&A	Hallucination; academic integrity; data privacy if student inputs are stored
Adaptive learning platforms	Adjusts content difficulty and sequencing based on individual student performance	Personalised practice, formative data, targeted intervention	Algorithmic bias in placement; data collection on minors; over-reliance displacing teacher judgment

AI writing tools	Assists with drafting, editing, grammar, style, and paraphrasing	Writing support, accessibility, ELL assistance	Academic integrity; dependency reducing skill development; grammar tools not flagging conceptual errors
AI image & media generators	Creates images, audio, or video from text prompts	Visual content creation, creative projects, accessibility (text-to-speech)	Copyright ambiguity; inappropriate output risk; deepfake potential; data training on user inputs
AI detection tools	Claims to identify AI-generated text in student submissions	Supporting academic integrity processes	High false-positive rates; bias against ELL students; not admissible as sole evidence; creates adversarial dynamic
Integrated AI in existing tools	AI features built into platforms teachers already use (Google Workspace, Microsoft 365, Canva, etc.)	Contextually embedded assistance within familiar workflows	Scope creep: features enabled by default without informed consent; data sharing between platforms

Section 2: Tool Profiles by Category

HOW TO READ THIS SECTION	Each table shows: Tool name What it does Best for in schools Watch for Cost (F=Free tier, P=Paid, F/P=Both) Data flag (✓=generally lower risk for school use, ⚠=verify privacy terms before any student use). These are not endorsements. Verify current terms before use.
--------------------------	--

AGeneral-Purpose Generative Text AI

The core category: chatbots and writing assistants that generate text from prompts

Tool	What it does	Best for	Watch for	Cost	Data
ChatGPT (OpenAI)	Conversational AI generating text, code, analysis, and images	Lesson planning, explanation, feedback drafts, teacher productivity	Default data used for training unless opted out; student use requires account (age 13+)	F/P	⚠
Claude (Anthropic)	Conversational AI with strong reasoning and longer context window	Complex explanation, feedback on extended writing, policy drafts	Privacy policy should be verified; no persistent memory by default	F/P	✓
Microsoft Copilot	AI integrated into Microsoft 365 — Word, Teams, Outlook, and more	Teachers using Microsoft 365 for lesson drafts, emails, meeting notes	Enterprise vs consumer versions have very different data terms; confirm	F/P	✓

			which version your school has		
Google Gemini	AI integrated into Google Workspace — Docs, Slides, Gmail, and more	Teachers using Google Workspace; summarising, drafting within existing tools	Workspace for Education has stronger privacy terms than consumer accounts; confirm school account type	F/P	✓
Khanmigo (Khan Academy)	Education-specific AI tutor and teacher assistant built on GPT	Direct student use for tutoring; Socratic prompting rather than giving answers	Requires Khan Academy account; designed for students 13+; school account needed for under-13	F/P	✓

B Adaptive Learning Platforms

AI-powered platforms that personalise practice and provide formative data

Tool	What it does	Best for	Watch for	Cost	Data
DreamBox (HMH)	Adaptive maths platform for K–8 that adjusts in real time	Maths differentiation and targeted practice	Collects significant student performance data; check state student privacy agreements	P	⚠
IXL Learning	Adaptive practice across maths, ELA, science and social studies K–12	Skill gap identification, targeted practice, teacher progress reports	Extensive data collection on student performance patterns	P	⚠
Formative	AI-assisted formative assessment with real-time response visibility	Formative checks, exit tickets, class-wide insight during lessons	Free tier limited; data stored on US servers — verify for non-US schools	F/P	✓
Eduaide.Ai	AI lesson planning assistant and resource generator for teachers	Generating lesson plans, rubrics, differentiated materials quickly	Teacher-facing productivity tool; less student-facing risk	F/P	✓

C AI Writing & Literacy Tools

Tools that assist with the writing process, grammar, and language

Tool	What it does	Best for	Watch for	Cost	Data
------	--------------	----------	-----------	------	------

Grammarly	Grammar, style, and clarity suggestions with AI-powered tone detection	Writing improvement, ELL support, professional writing for teachers	Free version stores text; Education version has better privacy terms; not suitable for unsupported assignments	F/P	⚠
Quillbot	AI paraphrasing, summarising, and grammar checking	ELL support, summarisation practice, paraphrasing exploration	High misuse potential for academic dishonesty; discuss explicitly with students before any school use	F/P	⚠
Diffit	Generates levelled readings and differentiated materials from any topic or text	Differentiation, accessible versions of complex texts, ELL support	Education-specific tool with reasonable privacy posture; verify before student use	F/P	✓
MagicSchool.ai	Suite of AI tools for teacher productivity: rubrics, lesson plans, reports	Reducing teacher administrative load; not designed for direct student use	Data policy evolving; teacher-facing use lower risk than student-facing	F/P	✓

D**AI Image & Media Generation***Tools that create visual, audio, or video content from prompts*

Tool	What it does	Best for	Watch for	Cost	Data
Canva AI (Magic Studio)	AI image generation, text-to-design, and presentation AI within Canva	Visual projects, presentations, creative assignments	School accounts available; copyright on AI-generated images ambiguous — school policy needed	F/P	✓
Adobe Firefly	AI image generation trained on licensed content (lower copyright risk than competitors)	Visual content creation where copyright clarity matters	Paid subscription; aimed at older students and professional use	P	✓
Suno / Udio	AI music generation from text prompts	Music projects, audio for presentations, creative exploration	Copyright of AI-generated audio unresolved; not school-vetted products — verify before any student use	F/P	⚠

Descript / Clip Champ	AI-assisted video and audio editing	Media production classes, student video projects, accessibility (captions)	Confirm school account vs consumer account terms; caution with student voice/image data	F/P	⚠
------------------------------	-------------------------------------	--	---	------------	---

E**AI Detection Tools***Tools that claim to identify AI-generated text in student submissions*

Tool	What it does	Best for	Watch for	Cost	Data
Turnitin AI Detection	AI detection integrated into Turnitin's plagiarism checking workflow	Part of an existing integrity process, not as sole evidence	Published false-positive rates; ELL student bias documented; one signal among many, never sole evidence	P	⚠
GPTZero	Standalone AI text detector designed specifically for education	Quick check as one signal in an investigation	Not reliable enough to be used as proof; creates adversarial student-teacher dynamic if over-relied upon	F/P	⚠
Copyleaks	Plagiarism and AI detection with LMS integration	Schools with existing Copyleaks contracts	Similar false-positive concerns apply; verify data residency for student submissions	P	⚠

IMPORTANT

AI detection tools should never be used as sole evidence of academic dishonesty. All major vendors publish false-positive rates. ELL students and students who write in a direct, clear style are disproportionately flagged. Use detection tools as one signal in a conversation, not as proof. See Module 3.1 for the full incident reporting process.

F**AI Features in Existing School Platforms***AI capabilities built into tools schools already use*

Tool	What it does	Best for	Watch for	Cost	Data
Microsoft 365 Copilot (EDU)	AI across Word, PowerPoint, Excel, Teams, and OneNote	Schools on Microsoft EDU licensing; teacher productivity	Separate licensing from standard M365; data protection	P	✓

			under Microsoft EDU terms generally strong		
Google Workspace AI (EDU)	Gemini AI features across Docs, Slides, Sheets, Meet, and Gmail	Schools on Google Workspace for Education; embedded teacher tools	Workspace for Education has stronger privacy terms than Gmail — confirm school account type before student use	F/P	✓
Canvas + AI features	AI-assisted assignment feedback and quiz generation in Canvas LMS	Schools using Canvas; reducing teacher grading load	Review Canvas AI addendum; student data processed by AI — check institutional data agreement	P	⚠
Schoology / PowerSchool AI	Analytics and predictive features in student information systems	District-level early warning systems, attendance/performance patterns	Most sensitive data type: grades, behaviour, demographics used in AI models — highest vetting standard required	P	⚠

Section 3: Eight Red Flags — When to Stop and Walk Away

These eight patterns indicate that a tool should not be used with students or student data until the concern is specifically resolved. They are ordered from most to least immediately disqualifying. Any one of them is sufficient reason to pause.

⚠ No accessible privacy policy, or one that contains student data

Every tool used in a school context must have a publicly accessible privacy policy. If you cannot find one in under two minutes, the tool should not be used with students. If the privacy policy mentions using user inputs to train models and does not explicitly exempt student data, that tool is collecting and potentially sharing student data without proper consent.

Example: A free AI quiz generator has no privacy policy link. Avoid entirely until resolved.

⚠ Claims that AI-generated scores or predictions are objective or unbiased

No AI system is unbiased. Any vendor that claims their tool produces objective assessments, unbiased predictions, or neutral recommendations should be treated with serious scepticism. Bias in AI is structural — it reflects training data, which reflects the world that data describes. Tools making objectivity claims are either poorly built or misleadingly marketed.

Example: *An adaptive assessment platform claims its AI ‘removes human bias from grading.’ This is not possible and not true.*

⚠ Tool requires students under 13 to create personal accounts

COPPA (Children's Online Privacy Protection Act) prohibits commercial websites from collecting personal information from children under 13 without verifiable parental consent. Many AI tools have a minimum age of 13 or 18. Any tool requiring students under 13 to create accounts with a name, email, or other personal information requires either parental consent processes or should not be used with that age group.

Example: *An AI writing tool requires a personal email address to register and has no school account option. Not suitable for K–5 use.*

⚠ Data is stored outside your jurisdiction without explicit consent

Student data stored on servers in different countries may be subject to different legal protections. Schools in the EU must comply with GDPR; many US states have additional student privacy laws. If a tool's privacy policy does not specify where data is stored, or specifies a jurisdiction with weaker protections, this requires legal review before school adoption.

Example: *A UK school identifies that a free EdTech tool stores student data on US servers only, with no EU data residency option.*

⚠ Tool uses AI detection as definitive proof of academic dishonesty

No current AI detection tool is reliable enough to serve as sole evidence of academic dishonesty. All major tools have documented false-positive rates and have shown disproportionate flagging of ELL students. Any school process or vendor tool that presents AI detection output as definitive evidence — rather than as one signal among many — is not following current best practice.

Example: *A school's integrity policy states that a Turnitin AI score above 20% constitutes proof of AI use, triggering automatic sanctions.*

⚠ Vendor makes therapeutic or emotional support claims without clinical backing

AI tools that position themselves as providing mental health support, emotional check-ins, or counselling for students are in a high-risk category. AI is not a licensed mental health professional, cannot provide crisis support reliably, and can cause harm if it gives inappropriate responses to distressed students. These tools require an extremely high bar of evidence and institutional oversight before any school adoption.

Example: *A free app offers ‘AI-powered emotional check-ins’ for students and suggests coping strategies based on their responses.*

⚠ The tool uses student performance data to sell or market to students or families

Some platforms use student performance and engagement data to target students or families with upsells, premium content, or third-party advertising. This is incompatible with educational use regardless of whether it is technically legal under the platform's terms.

Example: *An adaptive learning platform's terms allow sharing ‘aggregate and de-identified data with advertising partners’ — a common formulation that is not as protective as it sounds.*

⚠ No clear human oversight or appeal process in AI-generated decisions

Any AI system that makes or contributes to decisions affecting students — placement, grading, behavioural flags, resource allocation — must have a documented human oversight process and a clear way for students and families to appeal decisions. AI that makes high-stakes decisions without human review violates both ethical best practice and, in some jurisdictions, legal requirements.

Example: *A student information system uses AI to flag students as ‘at risk of dropout’ and automatically triggers interventions without teacher review.*

Section 4: The SAFE Evaluation Framework

When encountering any new or unfamiliar AI tool, apply the SAFE framework before adopting, recommending, or using it with students. Each letter represents a dimension of evaluation. All four dimensions must pass before a tool is suitable for school adoption.

S — Student data

The most critical dimension. Three questions: (1) Does the tool collect, store, or process any student data? (2) If yes, is there a school or district-level data processing agreement (DPA) in place? (3) Does the tool’s privacy policy explicitly protect student data from being used for AI training, advertising, or third-party sharing?

If the answer to question 1 is yes and question 2 or 3 is no, the tool should not be used with students until a DPA is in place.

Fail condition: *Fail condition: tool stores student inputs with no DPA and default terms allow training use.*

A — Age appropriateness

Content, design, and account requirements must be appropriate for the intended student age group. Check: minimum age in the terms of service; whether the tool requires personal account creation; whether content filters are in place for generative tools; and whether the tool’s pedagogical design matches the developmental stage of your students.

Distinct from data compliance: a tool can be COPPA-compliant but still developmentally inappropriate for a specific age group.

Fail condition: *Fail condition: tool requires students under 13 to create personal email accounts.*

F — Fit for purpose

Does the tool actually do what you need it to do, done well? Marketing language in the EdTech space is frequently aspirational rather than evidence-based. Questions to ask: Is there independent evidence of effectiveness (not just vendor case studies)? Can the teacher maintain meaningful oversight of what the tool does? Does the tool support learning or replace it? What happens to student learning if the tool is removed?

A tool that works well commercially may not serve your specific pedagogical goals.

Fail condition: *Fail condition: vendor claims ‘proven to raise test scores by 30%’ with no peer-reviewed evidence.*

E — Equity and bias

Does the tool perform consistently across different student groups? Key questions: Has the vendor published any equity analysis of tool performance? Are there documented performance gaps by language background, race, gender, or disability status? What happens when the tool makes an error — who is most affected? Is there a human override available?

Equity evaluation is often the hardest to conduct without vendor transparency, but it is not optional for any tool used in formative or summative assessment contexts.

Fail condition: *Fail condition: vendor has not conducted or published any equity analysis of adaptive scoring.*

KEY The SAFE framework is a minimum bar, not a maximum. A tool can pass all four dimensions and still not be the right choice for your school. Add your own contextual criteria — curriculum alignment, teacher comfort, infrastructure requirements, cost sustainability — after the four SAFE dimensions are cleared.

Section 5: AI Tool Vetting Checklist

Use this checklist when formally evaluating a new AI tool for school or classroom adoption. Complete one checklist per tool. A single ‘No’ in the Student Data section should halt the process until resolved. Multiple ‘Unsure’ responses indicate that the tool requires further investigation before any decision is made.

Tool name: _____ Evaluator: _____ Date: _____
Proposed use: _____

Question	Response	Notes / Evidence
S — Student Data		
Is there a publicly accessible privacy policy?	☐ Yes ☐ No ☐ Unsure	Notes:
Does the school / district have a signed Data Processing Agreement (DPA) with this vendor?	☐ Yes ☐ No ☐ Unsure	Notes:
Does the privacy policy explicitly state student data is NOT used for AI model training?	☐ Yes ☐ No ☐ Unsure	Notes:
Does the privacy policy state student data is NOT shared with third parties for advertising?	☐ Yes ☐ No ☐ Unsure	Notes:
Is data stored in a jurisdiction with adequate legal protection for our students?	☐ Yes ☐ No ☐ Unsure	Notes:
Is there a documented process for data deletion upon request?	☐ Yes ☐ No ☐ Unsure	Notes:

A — Age Appropriateness		
What is the minimum age requirement stated in the Terms of Service?	☐ Yes ☐ No ☐ Unsure	Notes:
Do students need to create personal accounts? If yes, what information is required?	☐ Yes ☐ No ☐ Unsure	Notes:
Does the tool have content filters appropriate to the intended age group?	☐ Yes ☐ No ☐ Unsure	Notes:
Is there a school / district account option that manages consent centrally?	☐ Yes ☐ No ☐ Unsure	Notes:
F — Fit for Purpose		
Is there independent (non-vendor) evidence of effectiveness for this tool?	☐ Yes ☐ No ☐ Unsure	Notes:
Can the teacher maintain meaningful oversight of what the tool does and shows students?	☐ Yes ☐ No ☐ Unsure	Notes:
Does the tool support and develop student learning, or replace student thinking?	☐ Yes ☐ No ☐ Unsure	Notes:
Is the tool technically accessible on school devices and networks?	☐ Yes ☐ No ☐ Unsure	Notes:
E — Equity and Bias		
Has the vendor published any equity or fairness analysis of tool performance?	☐ Yes ☐ No ☐ Unsure	Notes:
Are there documented performance gaps by race, language background, gender, or disability?	☐ Yes ☐ No ☐ Unsure	Notes:
Is there a human override or appeal process for AI-generated decisions?	☐ Yes ☐ No ☐ Unsure	Notes:
Are there known accessibility features for students with disabilities?	☐ Yes ☐ No ☐ Unsure	Notes:
Overall recommendation	☐ Approve ☐ Approve with conditions ☐ Further investigation needed ☐ Decline	
Conditions or concerns to resolve before approval:	_____	

Section 6: How to Read a Privacy Policy in Under 5 Minutes

Privacy policies are long, technical, and written for lawyers — but the information you need as an educator can be found with five specific searches. You do not need to read the whole document.

Search 1: Find ‘training’ or ‘improve our services’

Look for any clause stating that user inputs, data, or interactions are used to train, improve, or develop AI models. If this applies to student data without explicit exemption, it is a red flag. Many tools exempt ‘Enterprise’ or ‘Education’ accounts from training use — confirm which account type your school has.

Search 2: Find ‘third parties’ or ‘share’

Identify what data is shared with third parties and for what purposes. ‘Aggregate and de-identified data’ is a common formulation that sounds protective but can still allow identification in some cases. Look for whether advertising partners are listed.

Search 3: Find ‘children’ or ‘COPPA’ or ‘FERPA’

If the tool is used with students, these terms should appear. Their absence from a privacy policy for an educational tool is itself a red flag. The presence of COPPA or FERPA language does not guarantee compliance — but absence strongly suggests the vendor has not considered the school use case.

Search 4: Find ‘delete’ or ‘retention’

How long is student data kept? Is there a process for requesting deletion? FERPA requires that schools maintain control over student education records. If a vendor retains data indefinitely or has no deletion process, they are creating a liability for the school.

Search 5: Find ‘breach’ or ‘security incident’

What is the vendor’s process when data is breached? Schools typically have notification obligations to families and regulators. If the vendor’s policy does not include breach notification timelines, you cannot meet your own obligations.

TIP Many school districts negotiate Data Processing Agreements (DPAs) with commonly used vendors and publish these on their district website. Before independently vetting a tool, check whether your district already has a DPA in place. Common Student Privacy Consortium (studentprivacypact.org) also maintains a searchable database of vendor privacy commitments.

Facilitated Session Agenda (30 minutes)

Time	Segment	Activity	Format	Materials
0–3 min	Opening poll	Quick show of hands: which tools have you used in the past month? Establishes baseline and shared reference points.	Poll	Slides
3–10 min	The six categories	Facilitator walks through the category table. For each category, one participant names a tool they have encountered.	Input + share	Category table
10–18 min	Tool profile deep dive	In pairs, participants choose one tool category they know least about and read the profile table together. They write one question per tool.	Pairs	Tool profile tables
18–24 min	Red flag recognition	Facilitator presents three anonymised real-world tool descriptions. Groups identify which red flags apply.	Small groups	Red flag cards
24–28 min	SAFE framework drill	Participants apply the SAFE framework to a hypothetical new tool (facilitator reads the description). Groups reach a recommendation in 4 minutes.	Groups	SAFE summary card

28–30 min	Wrap-up	Vetting checklist orientation. Next steps: complete a checklist for one tool you currently use.	Facilitator	Vetting checklist
-----------	---------	--	-------------	-------------------

Activities — Facilitator Guide

A1 Category Sort

10 minutes

Participants receive a set of 12 AI tool description cards (no tool names, just what each does) and sort them into the six categories. The exercise builds pattern recognition and surfaces misconceptions about what different tools actually do — particularly the difference between generative AI and adaptive platforms, and between AI writing tools and AI detection tools.

1. Prepare 12 cards: two tools per category with their features described but not named.
2. Small groups (3–4 people) sort all 12 cards into the six categories. 6 minutes.
3. Groups compare sorts with each other. Where do they disagree? The most interesting disagreements are usually: AI writing tools vs generative AI (genuine grey area), and adaptive platforms vs integrated AI (depends on implementation).
4. Facilitator debrief: reveal the tool names. Ask 'Did knowing the name change your categorisation? Why or why not?' The goal is to reinforce that category belongs to mechanism, not brand.

A2 Red Flag Hunt

8 minutes

Groups receive three anonymised real-world tool descriptions taken from actual privacy policies and product pages (with vendor names removed). They identify which of the eight red flags apply. At least one description contains a disqualifying red flag; at least one is genuinely borderline. The exercise develops evaluation instincts rather than rule-following.

5. Distribute one tool description card per group (three different cards in rotation).
6. Groups read the description and identify: which red flag(s) apply? Is this a hard stop or a 'verify before proceeding'? (4 min)
7. Each group shares their finding. Facilitator adds nuance: some red flags are absolute; others depend on whether a DPA is in place or whether the tool is used teacher-facing vs student-facing.
8. Key message: 'Getting comfortable with this discomfort is part of the job now. You do not need to be a lawyer — you need to know when to call one.'

A3 SAFE Framework: Quick Evaluation

6 minutes

The facilitator reads a description of a hypothetical new AI tool — compelling marketing, some genuine red flags embedded — and groups have four minutes to produce a SAFE framework verdict: Approve, Approve with conditions, Further investigation, or Decline. This simulates the real-world decision a school leader or teacher might face when a vendor pitches a tool.

9. Facilitator reads the tool description aloud twice at normal pace. (1 min)
10. Groups discuss using the four SAFE dimensions and reach a verdict with one sentence of justification. (4 min)

11. Groups share verdicts. If they disagree, explore why: which dimension did different groups weight most heavily?
12. Facilitator reveals the 'answer' (prepared in advance): which real red flags were embedded in the description, and what the appropriate verdict was. The goal is not to catch participants out — it is to calibrate their evaluation instincts.

Participant Handouts — Content

Handout 1: Tool Categories Quick Reference Card

FORMAT Designed as a double-sided laminated A5 card. Front: the six category descriptions. Back: the eight red flags condensed to one line each.

Category	What it does	Examples	Primary risk
A Generative text AI	Generates text from prompts	ChatGPT, Claude, Copilot, Khanmigo	Hallucination; academic integrity; student data privacy
B Adaptive learning	Personalises content to individual performance	DreamBox, IXL, Formative	Data collection; algorithmic bias in placement
C AI writing tools	Supports drafting, editing, grammar	Grammarly, Quillbot, Diffit	Academic integrity; dependency; paraphrasing misuse
D Image & media AI	Creates images, audio, video from text	Canva AI, Adobe Firefly, Suno	Copyright ambiguity; inappropriate output; deepfake risk
E AI detection	Claims to identify AI-generated text	Turnitin AI, GPTZero	False positives; ELL bias; cannot be sole evidence
F Integrated AI	AI built into existing school platforms	M365 Copilot, Google Gemini, Canvas AI	Scope creep; default-enabled data sharing

Handout 2: SAFE Framework Summary Card

Letter	What to assess	Key question	Fail condition
S Student data	Data collection, storage, training use, third-party sharing, DPA existence	<i>What happens to the data students put into this tool?</i>	No DPA; data used for training; no deletion process
A Age appropriateness	Minimum age, account requirements, content filters, developmental fit	<i>Is this appropriate for my specific students' age group?</i>	Under-13 students required to create personal accounts
F Fit for purpose	Independent evidence, teacher oversight, learning support vs replacement	<i>Does this actually help my students learn, with evidence?</i>	Vendor claims unsupported by

			independent evidence
E Equity and bias	Equity analysis, performance gaps, human override, accessibility	Has the vendor shown this works fairly for all student groups?	No equity analysis published; no human override for AI decisions

Facilitation Notes & Common Scenarios

Managing Tool Currency

The specific tools in this module will change. New tools will emerge; existing tools will update their pricing, features, or terms. When facilitating this module, spend a few minutes before each delivery checking whether the cost, availability, or data terms of any featured tool have changed significantly. The frameworks and red flags are durable; the specific tool details are not.

Anticipated Participant Reactions

“I already use several of these tools and didn’t know about the data issues.”

This is the most common reaction, and it is a productive one. Acknowledge it without alarm: ‘Many educators have been in that position. The goal today is not to create anxiety about past decisions but to build the habits that make future decisions more informed.’ The vetting checklist is the practical takeaway: start using it going forward, even if past adoptions weren’t vetted.

“Our IT department handles all of this — why do I need to know it?”

IT handles institutional procurement, but teachers frequently recommend or adopt tools independently — particularly free tools that don’t require IT approval. Every time a teacher directs students to a free web tool or recommends an app, that is a data decision being made without IT oversight. This module builds the judgment teachers need for the decisions that happen at the classroom level.

“These tools are all banned at our school anyway.”

Blanket bans are one policy response, but they create their own risks: students and teachers use tools through personal devices, outside school networks, where the school has even less oversight. Understanding the landscape is valuable whether your school uses these tools or restricts them. And policies evolve — being informed when that happens is more useful than being caught unprepared.

“AI detection tools caught a student last term. Are you saying I was wrong to use it?”

Handle this carefully: acknowledge that AI detection can be one signal in an investigation, and that a teacher who used it alongside other evidence was doing their job. The concern is when detection output is treated as proof rather than as a signal. Ask: ‘Was the detection flag the only evidence, or one of several things you noticed? Did you have a conversation with the student before any formal process?’ The Module 3.1 incident reporting process covers this in detail.

Related Modules in the PD Series

Recommended before this session

- 1.1 AI Self-Literacy Assessment

Recommended after this session

- 2.3 AI Prompt Literacy for Teachers

- 1.2 How AI Actually Works — A Practical Primer

- 3.1 AI Incident Identification & Reporting
- 3.3 Student Data Privacy in the Age of AI
- 4.3 AI Vendor Evaluation & Procurement
- 4.1 Building Your School AI Policy

TIP

This module pairs especially well with Module 4.3 (AI Vendor Evaluation & Procurement) for school administrators and IT staff. Teachers who complete Module 1.3 develop individual evaluation skills; Module 4.3 builds the institutional process for formal procurement. Ideally both are completed in the same school year by any team making AI adoption decisions.