

PROFESSIONAL DEVELOPMENT MODULE

How AI Actually Works

A Practical Primer for Educators

The essential concepts behind AI — explained without code, without hype, and without the jargon. What educators need to know to teach, lead, and decide with confidence in an AI-present school.

Module 1.2 • Track 1: AI Foundations & Literacy • xvalidateai.com

Audience K–12 Teachers & Leaders	Format Micro-module	Duration 25–30 minutes	Level Intermediate	Delivery Self-paced or facilitated
--	-------------------------------	----------------------------------	------------------------------	--

Module Overview

You do not need to understand how an engine works to drive a car well. But you do need to know what fuel it takes, roughly how it behaves in different conditions, and what the dashboard warning lights mean. This module gives educators the equivalent knowledge about AI — not a technical deep dive, but the conceptual literacy that lets you make good decisions, spot problems, and talk confidently with students, parents, and administrators.

This module is designed as a 25–30 minute self-paced experience, or can be facilitated in a group setting as a focused micro-session. It combines short readings, worked analogies, a myth-busting reference sheet, and a practical vocabulary guide. No prior AI knowledge is assumed beyond having used the tools at least once.

WHO THIS IS FOR

This module is the right starting point if you feel like you understand AI at a surface level from using ChatGPT, Copilot, or similar tools, but would not feel confident explaining how it works to a student or parent, or explaining why it sometimes gets things wrong.

Learning Objectives

By the end of this module, participants will be able to:

- | | |
|---|--|
| 1 | Explain in plain language what an AI language model is and what it is actually doing when it generates a response — without relying on technical jargon |
| 2 | Distinguish between different types of AI (narrow AI, generative AI, machine learning) and explain why these distinctions matter in an educational context |

3	Identify and correct the five most common misconceptions about AI held by students and adults
4	Use the seven key AI vocabulary terms accurately in professional conversations with colleagues, students, and families
5	Explain three structural reasons why AI makes mistakes — and connect these to practical implications for how teachers and students should use AI outputs
6	Articulate a personal, evidence-based answer to the question: 'Is AI intelligent?'

Part 1: What AI Actually Is

The one-sentence version

A modern AI language model is a very large statistical system trained on enormous amounts of text that has learned to predict what words are most likely to come next, given what has come before — and through doing this at extraordinary scale and depth, it produces outputs that often appear intelligent, creative, and knowledgeable.

ANCHOR

The key word in that sentence is 'predict.' AI is not retrieving facts from a database. It is not reasoning the way a human reasons. It is completing patterns — at a scale and sophistication that produces results that can look remarkably like thinking.

Three analogies — each one useful, each one limited

No single analogy captures what AI is doing. The three below each illuminate something true and obscure something important. Understanding all three — including their limits — gives a more accurate picture than any one alone.

Analogy: Autocomplete at extraordinary scale

What the analogy captures

Just like your phone predicts the next word when you type, AI predicts what comes next — but trained on billions of documents, doing it for entire paragraphs, essays, code, and conversation. This is why it is fluent and coherent: it has seen enormous numbers of examples of how language flows.

What's technically happening

The model is generating each token (roughly each word) one at a time, sampling from a probability distribution over all possible next tokens based on everything that came before in the conversation. The 'knowledge' is encoded in hundreds of billions of numerical weights, not a lookup table.

Where the analogy breaks down:

Analogy: A very well-read graduate student with no memory

What the analogy captures

AI has read more than any human could in a lifetime: articles, books, code, conversations. It can synthesise, explain, argue, and create

What's technically happening

Each conversation starts completely fresh. The model has no memory of previous conversations (unless explicitly provided). It

across a vast range of topics. Like a knowledgeable graduate student, it can engage substantively with almost any subject.	has no ongoing experience of the world. Its 'knowledge' has a cutoff date. And unlike a graduate student, it has no stake in being right — it is optimising for plausible output, not truth.
Where the analogy breaks down:	

Analogy: A distillation of human expression	
What the analogy captures AI has been trained on the accumulated written output of human civilisation: scientific papers, novels, arguments, poetry, journalism, conversations. In a sense, it reflects back a compressed version of how humanity has expressed ideas — which is why it can produce content that feels recognisably human.	What's technically happening The training data is not a neutral or representative sample of human thought. It skews heavily toward English and Western sources, toward text that was digitised and publicly available, and toward whatever was abundant on the internet. AI reflects not 'humanity,' but a particular filtered slice of it.
Where the analogy breaks down:	

The difference between types of AI

The word 'AI' covers a wide range of different technologies. Understanding the distinctions helps educators avoid both over-claiming and under-claiming about what AI can and cannot do.

Concept	Plain-language definition	Why it matters for educators
Narrow AI	AI designed to do one specific task very well — like recognising a face, recommending a video, or translating a sentence. It cannot generalise beyond its task.	The AI that flags plagiarism, the speech-to-text tool, the adaptive maths software — all narrow AI. Being good at one thing says nothing about being good at anything else.
Machine learning	A method of building AI systems by training them on examples rather than programming rules. The system learns patterns from data rather than following explicit instructions.	Most modern AI tools, including generative AI, use machine learning. The 'learning' happened during training — the system is not learning in real time during your conversation.
Generative AI	AI that produces new content — text, images, code, audio, video — rather than just classifying or retrieving. ChatGPT, Claude, Copilot, Gemini, Dall-E, and Midjourney are generative AI.	This is the category most relevant to student use. Generative AI is powerful, creative, and error-prone in ways that narrow AI is not.
Large Language Model (LLM)	The specific type of generative AI that processes and produces text. 'Large' refers to the scale of	When students use ChatGPT or Claude to write, explain, or answer questions, they are using an LLM.

	training data and model parameters — billions to trillions.	Understanding this category is the most practically relevant for classroom decisions.
General AI (AGI)	A hypothetical future AI that could perform any intellectual task a human can. Does not currently exist.	AGI is often what people imagine when they worry about AI — or when they over-attribute capability to current tools. Current AI is powerful but narrow compared to the concept of AGI.

Part 2: Why AI Makes Mistakes — and Why It Matters

Understanding why AI makes mistakes is not a technical detail — it is a practical necessity for anyone who uses AI tools or teaches students who use them. There are three structural reasons, each with direct educational implications.

Reason 1: AI does not know what it does not know

AI language models generate outputs based on patterns in their training data. When they encounter a question that is outside that data, or where the training data was incorrect or sparse, they do not reliably produce uncertainty. They often produce a plausible-sounding answer at the same confidence level as a correct one. This is called hallucination — the generation of false information with apparent fluency.

What hallucination looks like

A fabricated academic citation that looks real but does not exist
A confident statement about an event that never happened
A plausible but incorrect historical date or statistic
A medical or legal claim that is close to true but critically wrong in detail

Educational implication

Every specific fact, statistic, citation, or date produced by AI must be independently verified before it is used in student work or teacher materials. This is not an occasional precaution — it is a standing practice. The fluency of the output is not a signal of its accuracy.

Reason 2: AI’s knowledge has a cutoff date

AI language models are trained on data collected up to a specific point in time — their ‘knowledge cutoff.’ Events, research, policies, and discoveries after that date are not in the model’s training data. Some models have tools that allow them to search the web in real time, but the base model itself has no awareness of anything that happened after its cutoff.

IMPLICATION

This matters most for: current events and news, recent scientific findings, updated laws or policies, and new or recently released educational resources. Asking AI about ‘the latest research on reading instruction’ or ‘current state assessment requirements’ is particularly risky without real-time search enabled.

Reason 3: AI reflects the biases of its training data

Because AI is trained on human-generated text, it absorbs the patterns, assumptions, and biases present in that text. This includes cultural biases, historical biases embedded in how events were described, representation gaps (who gets written about, in what contexts, by whom), and the particular skews of the internet — which does not represent all people, languages, or perspectives equally.

Bias in AI outputs is not always obvious. It can appear as: the names used in examples defaulting to certain cultural backgrounds, descriptions of professions associating with particular genders, historical narratives told from particular perspectives, or advice that assumes a particular cultural context as the default.

IMPLICATION

For educators, this means actively evaluating AI outputs for representational fairness, particularly when using AI to generate examples, scenarios, or explanatory text that will be shown to students. A diverse group of students deserves to see themselves reflected accurately in classroom materials.

Part 3: Five Myths Worth Busting

These five misconceptions are widespread among students, parents, and educators. Each one has practical consequences for how AI gets used in school. Knowing the reality behind each myth makes you better equipped to have productive conversations about AI with your school community.

MYTH → AI is searching the internet for answers

REALITY: Most AI language models generate responses from patterns in their training data, not by searching the web in real time. Some models have optional search tools, but by default they are producing outputs from internal weights, not live retrieval. When you ask ChatGPT or Claude a question, the response is generated, not looked up.

WHY IT MATTERS IN SCHOOL: Students who believe AI is searching the internet may assume its outputs are current, sourced, and verifiable in the way a Google result would be. They need to know that a fluent, confident AI answer is not the same as a retrieved fact.

MYTH → AI understands what it's saying

REALITY: This is one of the most genuinely contested questions in AI research, and the honest answer is: we do not know, and 'understanding' may not be the right frame. What we can say is that AI produces outputs as if it understands, through statistical pattern completion — but it has no experiences, intentions, or awareness of the world.

WHY IT MATTERS IN SCHOOL: When AI confidently explains a concept, students may assume the explanation is correct because it 'sounds like it knows.' AI produces fluent explanations of things it is wrong about just as readily as things it is right about. Fluency is not a signal of accuracy.

MYTH → AI is neutral and objective

REALITY: AI reflects the patterns of its training data — including all the biases, gaps, and cultural assumptions embedded in that data. It is not neutral. Neutrality would require either perfectly unbiased training data (which does not exist) or active debiasing (which is an imperfect, ongoing process). AI systems vary in how much debiasing work has been done.

WHY IT MATTERS IN SCHOOL: If students or teachers treat AI output as a neutral arbiter of fact or as a culturally neutral source of examples, they will fail to notice the biases embedded in those outputs. Critical evaluation of AI content needs to include asking: whose perspective is this reflecting?

MYTH → AI will replace teachers

REALITY: AI can automate some tasks that teachers currently do manually — generating initial drafts, producing practice questions, explaining concepts. But teaching is not primarily an information-delivery function. It involves relationships, professional judgment, emotional attunement, ethical decision-making, and responsiveness to individual human beings in a room — none of which AI can replicate.

WHY IT MATTERS IN SCHOOL: Fear of replacement can make teachers defensive about AI rather than strategic about it. The more useful framing: AI changes which parts of teaching take the most time and effort, and which skills matter most. The goal is to understand those changes and adapt — not to ignore them or be paralysed by them.

MYTH → AI is always getting better at everything

REALITY: AI capabilities are advancing rapidly in some dimensions and more slowly in others. AI can still make basic logical errors, fails at spatial reasoning tasks that children find easy, and struggles with tasks requiring common-sense understanding of the physical world. The improvements are real but uneven, and confident predictions about AI capability timelines have repeatedly been both wrong in each direction.

WHY IT MATTERS IN SCHOOL: Treating AI as infallible or as perpetually improving obscures the real picture: AI has specific, somewhat predictable failure modes that are worth understanding. Teaching students to probe for those failure modes is a valuable critical thinking skill.

Part 4: Is AI Intelligent? — A Question Worth Discussing

This is the question students ask most frequently, and it is worth having a considered answer. The honest response involves distinguishing between different meanings of ‘intelligent’ — and acknowledging genuine uncertainty.

Arguments that AI is intelligent

It performs tasks that, when done by humans, require intelligence — writing, reasoning, problem-solving, translation, explanation.

It can engage with novel problems it was not explicitly trained on.

Its performance on standardised tests (SAT, bar exam, medical licensing) meets or exceeds average human performance.

It generates explanations that reflect genuine conceptual relationships, not just surface word patterns.

Arguments that AI is not intelligent

It has no understanding, intentions, desires, experiences, or consciousness in any sense we can confirm.

Its ‘reasoning’ does not generalise the way human reasoning does — it breaks down in ways humans do not.

It cannot learn from errors in real time, update its beliefs based on new evidence, or set its own goals.

It does not know when it is wrong, and does not experience the dissonance of making a mistake.

SYNTHESIS

The most useful position for educators: AI is a powerful tool that performs some cognitively demanding tasks extremely well, fails at others in ways that surprise us, has no inner life that we can confirm, and raises genuinely novel philosophical questions about the nature of intelligence. Engaging with those questions — rather than defaulting to either ‘AI is just autocomplete’ or ‘AI is basically human’ — is both more accurate and more educationally valuable.

A note on anthropomorphism

AI tools are designed to communicate in natural language — which naturally produces language patterns that feel personal: ‘I think,’ ‘I’d be happy to,’ ‘I’m not sure about that.’ This is a design choice, not evidence of inner experience. Students often develop parasocial feelings toward AI chatbots and are at risk of attributing emotional states, intentions, and reliability to a system that has none of these in any stable sense. Educators can help by naming this dynamic explicitly and returning, regularly, to what the tool is actually doing.

Part 5: Essential Vocabulary Guide

These seven terms come up repeatedly in discussions about AI in education. Each definition below is written to be used directly with colleagues, students, or parents — plain, accurate, and free of unnecessary jargon.

Concept	Plain-language definition	Why it matters for educators
Training data	The collection of text (and other content) that an AI model learned from before it was released. Most large language models were trained on billions of documents from the internet, books, and other sources.	Training data explains why AI knows so much, why its knowledge has a cutoff date, why it reflects certain cultural biases, and why it may know more about some subjects than others.
Parameters / weights	The numerical values inside an AI model that store everything it learned during training. A large language model may have hundreds of billions of parameters. These are not facts stored in a database — they are patterns encoded as numbers.	Helps explain why AI cannot be ‘updated’ easily like a database, and why the model’s knowledge is probabilistic rather than look-up-based.
Token	The basic unit AI processes — roughly a word or part of a word. AI generates responses one token at a time, each based on what came before.	Useful for understanding why AI outputs flow grammatically even when factually wrong, and why very long inputs can cause quality to drop.
Prompt	The input you give to an AI — your question, instruction, or context. The quality and structure of the prompt has a large effect on the quality of the output.	The foundation of prompt literacy (Module 2.3). Understanding that the prompt is a design decision — not just a question — is the most

		practically useful conceptual shift for educators.
Hallucination	When an AI generates information that is false, fabricated, or unsupported — but stated with fluency and apparent confidence. The term comes from the idea of the model ‘perceiving’ something that is not there.	The most important term for educators to know and explain to students. AI hallucinations can look identical to correct information. Independent verification is always necessary for specific factual claims.
Context window	The amount of text an AI can ‘see’ at once in a single conversation. Earlier in a long conversation, things the AI was told may effectively be forgotten as new content pushes them out of the context window.	Explains why AI may seem to ‘forget’ instructions given early in a long conversation, and why providing key context at the start of a prompt is important.
Fine-tuning	A further training process that adjusts a model’s behaviour for a specific purpose — e.g., making it more helpful, safer, more focused on a particular domain, or better at following instructions.	Explains why different AI products based on the same underlying model can behave very differently — and why an AI tool marketed for education may behave differently from the general-purpose version of the same model.

Part 6: What This Means for Educators

The purpose of understanding how AI works is not academic — it is practical. The following implications follow directly from the concepts in this module.

VERIFY Check before you trust

Because AI hallucinate with the same fluency as accurate outputs, every specific factual claim — citation, statistic, date, name, policy detail — requires independent verification before it is used in teacher materials or shared with students. This is not occasional good practice; it is a standing requirement.

PROMPT The input shapes the output

Because AI produces responses based on what it is given, the structure and quality of a prompt directly determines the quality of the output. Understanding this shifts the challenge from ‘AI is unpredictable’ to ‘I can get better results by being more specific.’ Prompt literacy (Module 2.3) is a practical extension of this insight.

BIAS Look for whose voice is missing

Because AI reflects the biases of its training data, any AI-generated content used in the classroom should be evaluated not only for accuracy but for representational fairness. Whose perspective is centred? Whose experience is assumed to be universal? What cultural context is treated as the default?

TEACH

Students who understand how AI works are better equipped to use it well, evaluate its outputs critically, and make informed decisions about when to use it and when

Name the mechanism with students

not to. Sharing a simplified version of the concepts in this module — appropriate to their age — is one of the highest-leverage things a teacher can do.

LEAD

Inform policy with accurate understanding

School policies about AI — what is permitted, what is prohibited, how it is handled in assessment — should be grounded in an accurate understanding of what AI actually is and does. Policies based on misconceptions (e.g., that AI detection tools are reliable, or that AI use is always obvious) will be ineffective, inequitable, or both.

Self-Paced Module Structure

This module is designed for asynchronous delivery as well as facilitated use. The self-paced version breaks the content into four reading units, each followed by a short embedded reflection. Total time: 25–30 minutes.

Unit	Title	Core content	Reflection task	Time
1	What AI is	Parts 1–2: the one-sentence definition, three analogies, types of AI	<i>In one sentence, complete: ‘AI is... but it is not...’</i>	8 min
2	Why it goes wrong	Part 2: hallucination, knowledge cutoff, bias — with examples	<i>Name one time you encountered each type of AI error (or one you can imagine encountering)</i>	7 min
3	Myths and intelligence	Parts 3–4: myth-busting cards, the intelligence question	<i>Which myth do you hear most often? How will you respond to it now?</i>	7 min
4	Vocabulary and implications	Parts 5–6: seven terms, five practical implications	<i>Which implication is most relevant to your current school context?</i>	6 min

Facilitated Session Agenda (30 minutes)

Time	Segment	Activity	Format	Materials
0–4 min	Opening question	‘In one sentence, how would you explain to a parent what AI actually is?’ Pair share — no wrong answers, just surface the current baseline.	Pair share	Slides

4–10 min	What AI is	Facilitator delivers the one-sentence definition, walks through all three analogies with commentary on where each breaks down	Facilitator input	Slides
10–18 min	Why it goes wrong	Three reasons with real examples. Participants share one example each from their own experience — ‘an AI output that surprised you.’	Input + examples	Slides
18–24 min	Myth-busting pairs	Pairs receive two myth cards each. They read the myth and the reality, then discuss: ‘How would you explain this to a student?’	Pair work	Myth cards
24–28 min	Vocabulary speed-round	Facilitator reads the plain-language definition; participants name the term. Fast-paced, low-stakes. Correct any misconceptions.	Whole group	Vocab card
28–30 min	Implications & next steps	Which implication matters most to you right now? Name one thing you will do differently this week.	Quick share	Implications card

Activities — Facilitator Guide

A1 Three-Analogy Workshop

8 minutes

Participants work in small groups with all three analogies printed on cards. Each group evaluates the strengths and limits of their assigned analogy, then reports back. The goal is to develop a nuanced, multi-model understanding rather than defaulting to the first analogy that feels intuitive.

1. Distribute one analogy card per group of three participants. Each card has: the analogy, space for notes on what it captures well, and space for notes on where it breaks down.
2. Groups have 5 minutes to discuss both sides of their analogy.
3. Each group shares their ‘biggest limit’ — the place where their analogy most misleads (2 min).
4. Facilitator synthesis: ‘No single analogy is complete. The goal is not to pick the best one — it’s to hold all three simultaneously and notice what each illuminates.’

A2 Hallucination Hunt

10 minutes

Participants use a live AI tool (or pre-prepared printed examples if devices are unavailable) to test for hallucination. They ask the AI about a specific topic within their own subject area that they know well — and identify any inaccuracies in the output. This makes hallucination concrete rather than theoretical.

5. Participants open an AI tool (or receive the pre-prepared hallucination example sheet with three outputs from different subject areas).

6. They ask the AI a specific factual question from their own subject that they can verify personally. Suggested starters: ask for specific dates, names, or citations within a topic they know well.
7. Participants annotate the output: mark anything that is inaccurate, check any claims they are unsure about, and note whether the AI expressed any uncertainty about uncertain claims.
8. Pair discussion: 'If a student had submitted this output as their own research, what would you have noticed? What would you have missed?'
9. Debrief: the goal is not to conclude 'AI is unreliable' — it is to conclude 'AI outputs on factual claims need verification, and fluency is not a signal of accuracy.'

A3 Vocabulary Check-In — Plain Language Practice

5 minutes

A fast-paced whole-group activity to consolidate vocabulary. Facilitator reads a plain-language definition; participants name the term. Then one participant explains the educational significance of that term in one sentence. Keeps energy high and surfaces any remaining confusion.

10. Facilitator reads each definition from the vocabulary guide at normal speaking speed (no pauses for effect).
11. First person to name the term correctly — or collective call-out — earns the point if scoring is used.
12. After the term is named, a different participant explains: 'Why does this matter for teachers?' One sentence only.
13. Facilitator corrects any misattributions and reinforces the most important educational implication for any terms the group found difficult.

Participant Handouts — Content

Handout 1: The AI Vocabulary Card

FORMAT

Designed as a double-sided A5 laminated card. Front: the seven terms. Back: the three analogies with their limits. A quick reference for any professional conversation about AI.

Term	Plain-language definition
Training data	The collection of text and content an AI model learned from. It shapes what the AI knows, what it doesn't, and what biases it has absorbed.
Parameters	The billions of numerical values inside an AI that store everything it learned. Not facts in a database — patterns encoded as numbers.
Token	The basic unit AI processes — roughly a word or part of a word. AI generates one token at a time, each influenced by everything before it.
Prompt	The input you give to AI: your question, instruction, or context. Prompt quality directly shapes output quality.

Hallucination	When AI generates false information stated with confidence and fluency. It does not know it is wrong. Verification is always necessary for factual claims.
Context window	How much text AI can 'see' at once in a conversation. Information early in a long conversation may effectively be forgotten.
Fine-tuning	A further training process that adjusts a model's behaviour for a specific purpose. Why different AI products based on the same model behave differently.

Handout 2: The Five Myths Reference Card

A condensed single-page reference covering all five myths, their realities, and the educational implication of each. Designed to be shared with colleagues who have not completed this module.

Myth	Reality	Educational implication
AI searches the internet	Most AI generates from training data, not live web search	Outputs are not automatically current or sourced — verify specific facts
AI understands what it says	AI completes patterns; whether it 'understands' is genuinely uncertain	Fluency does not equal accuracy — never assume a clear explanation is a correct one
AI is neutral and objective	AI reflects biases embedded in training data	Evaluate AI content for representational fairness, especially in student-facing materials
AI will replace teachers	Teaching requires relationships, judgment, and human responsiveness AI cannot replicate	The goal is strategic adaptation, not fear — understand what changes, not whether you're at risk
AI always gets better at everything	Improvements are real but uneven; AI has predictable failure modes	Teaching students to probe AI limitations is a valuable critical thinking skill

Handout 3: Analogy Comparison Card

Analogy	What it captures well	Where it breaks down
Autocomplete at scale	Why AI outputs are fluent and coherent; why it 'knows' how language works	Risks making AI seem simpler than it is; autocomplete implies shallow pattern-matching only
Well-read grad student (no memory)	The breadth of knowledge; the ability to engage substantively; the need to start fresh each time	A human knows when they don't know; AI does not reliably express uncertainty about uncertain claims

Distillation of human expression	Why AI can produce recognisably human-sounding content; why it reflects human culture	Implies a neutral, representative sampling; the training data is neither neutral nor representative
---	---	---

Facilitation Notes

Managing the Depth vs. Time Tension

This module covers genuinely complex territory in 25–30 minutes. The risk is either going too shallow (participants leave with buzzwords but no real understanding) or too deep (the session runs over and participants feel overwhelmed). The key is anchoring every concept to a practical implication, and being comfortable saying ‘that is a deeper question than we have time for today — Module X goes further.’

Anticipated Participant Reactions

“I already knew most of this.”

Acknowledge and redirect to depth: ‘You knew the concepts — what we’re adding today is the precision to explain them accurately to others, and to recognise the specific failure modes that matter in school.’ Ask them: ‘Could you explain what hallucination means and why it matters to a Year 8 student in two sentences?’

“This is too technical for me to explain to students.”

The vocabulary guide and analogies are specifically designed to be used directly with students and parents. Invite participants to try the one-sentence definition: ‘AI predicts what comes next based on patterns in billions of texts — which is why it’s fluent but not always right.’ That is accessible to most secondary students and many primary students with a little scaffolding.

“Is AI going to get better and fix all of these problems?”

Honest answer: some of these limitations are structural (bias reflects training data; that data is human-generated) and will not be eliminated by scaling. Others (hallucination rates, knowledge cutoff freshness) are improving but not solved. We teach to current reality, not to a projected future that may or may not arrive on any particular timeline.

“Do I really need to understand this to just use AI tools?”

No — you can use a car without understanding combustion engines. But you need to know what fuel it takes, roughly how it behaves in different conditions, and what the warning lights mean. This module is the educational equivalent of that driving literacy, not mechanical engineering.

Related Modules in the PD Series

Start here first

- 1.1 AI Self-Literacy Assessment

Natural next steps

- 1.3 AI Tools Landscape for Educators
- 2.1 Designing AI-Resilient Assessments
- 2.2 Teaching Students to Use AI Responsibly
- 2.3 AI Prompt Literacy for Teachers
- 3.2 Bias, Fairness & AI in Education

TIP

This module is the conceptual foundation for everything else in the curriculum. It should precede any session that asks teachers to make policy decisions (Track 4) or teach students about AI (Track 2). A staff body that shares accurate conceptual understanding of AI makes far better collective decisions than one operating on a mix of myths and marketing.